

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: path=r"C:\Users\kasho\OneDrive\Documents\Data Science\flask\Salary_Data - S
df=pd.read_csv(path)
df.head()
```

Out[2]:

|   | YearsExperience | Salary |
|---|-----------------|--------|
| 0 | 1.1             | 39343  |
| 1 | 1.3             | 46205  |
| 2 | 1.5             | 37731  |
| 3 | 2.0             | 43525  |
| 4 | 2.2             | 39891  |

```
In [3]: df.shape
```

Out[3]: (30, 2)

```
In [4]: df.columns
```

Out[4]: Index(['YearsExperience', 'Salary'], dtype='object')

```
In [5]: df.isnull().sum()
```

Out[5]: YearsExperience 0  
Salary 0  
dtype: int64

```
In [6]: df.dtypes
```

Out[6]: YearsExperience float64  
Salary int64  
dtype: object

```
In [7]: X=df.drop('YearsExperience',axis=1)
y=df['Salary']
```

```
In [8]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test=train_test_split(X, # Input data
                                                y, # output data
                                                random_state=1234, # it s
                                                test_size=0.30)
```

```
In [9]: X_train.shape,X_test.shape
```

```
Out[9]: ((21, 1), (9, 1))
```

```
In [10]: y_train.shape,y_test.shape
```

```
Out[10]: ((21,), (9,))
```

```
In [11]: df.shape
```

```
Out[11]: (30, 2)
```

```
In [12]: X_train
```

```
Out[12]:
```

|    | Salary |
|----|--------|
| 13 | 57081  |
| 22 | 101302 |
| 24 | 109431 |
| 0  | 39343  |
| 2  | 37731  |
| 27 | 112635 |
| 26 | 116969 |
| 18 | 81363  |
| 5  | 56642  |
| 16 | 66029  |
| 25 | 105582 |
| 11 | 55794  |
| 9  | 57189  |
| 17 | 83088  |
| 29 | 121872 |
| 20 | 91738  |
| 12 | 56957  |
| 21 | 98273  |
| 6  | 60150  |
| 19 | 93940  |
| 15 | 67938  |

```
In [13]: y_train
```

```
Out[13]: 13      57081
          22     101302
          24     109431
           0      39343
           2      37731
          27     112635
          26     116969
          18      81363
           5      56642
          16     66029
          25     105582
          11      55794
           9      57189
          17     83088
          29     121872
          20      91738
          12      56957
          21      98273
           6      60150
          19      93940
          15      67938
          Name: Salary, dtype: int64
```

```
In [14]: X_test
```

```
Out[14]:
```

|    | Salary |
|----|--------|
| 7  | 54445  |
| 10 | 63218  |
| 4  | 39891  |
| 1  | 46205  |
| 28 | 122391 |
| 8  | 64445  |
| 3  | 43525  |
| 23 | 113812 |
| 14 | 61111  |

```
In [15]: y_test
```

```
Out[15]: 7      54445
          10     63218
           4      39891
           1      46205
          28     122391
           8      64445
           3      43525
          23     113812
          14      61111
          Name: Salary, dtype: int64
```

```
In [16]: X_train.ndim
# 1 dimension means 1 column only
# 2 dimension means 2 column only
# when you have only 1 column, the shape will not show the column
# (21,) it is only one column data having 21 observations
# (9,) it is one column data having 9 observation
# (30,2) it is 2 column data having 30 observation
# Reshape the data if you have only one column
```

Out[16]: 2

```
In [17]: from sklearn.linear_model import LinearRegression
LR=LinearRegression()
LR.fit(X_train,y_train)
```

Out[17]:

▼ LinearRegression

LinearRegression()

```
In [18]: # Model predictions happens X_test
y_predictions=LR.predict(X_test)
```

```
In [19]: y_predictions
```

Out[19]: array([ 54445., 63218., 39891., 46205., 122391., 64445., 43525.,  
113812., 61111.])

```
In [20]: y_test.shape,y_predictions.shape
```

Out[20]: ((9,), (9,))

```
In [21]: X_test
```

Out[21]:

|    | Salary |
|----|--------|
| 7  | 54445  |
| 10 | 63218  |
| 4  | 39891  |
| 1  | 46205  |
| 28 | 122391 |
| 8  | 64445  |
| 3  | 43525  |
| 23 | 113812 |
| 14 | 61111  |

```
In [22]: X_test.iloc[0] # series
# In order to pass a test sample to a model
# we need to pass a list of values
# or array of values
# tuple of values
X_test.iloc[0].values
```

```
Out[22]: array([54445], dtype=int64)
```

```
In [23]: LR.predict([X_test.iloc[0].values,
                     X_test.iloc[1].values])
```

```
Out[23]: array([54445., 63218.])
```

```
In [24]: ip1=[5]
LR.predict([ip1])
```

```
Out[24]: array([5.])
```

```
In [25]: X_test.shape,y_test.shape,y_predictions.shape
```

```
Out[25]: ((9, 1), (9,), (9,))
```

```
In [26]: test_data=X_test
test_data['y_actual']=y_test
test_data['y_predictions']=y_predictions
test_data
```

```
Out[26]:
```

|    | Salary | y_actual | y_predictions |
|----|--------|----------|---------------|
| 7  | 54445  | 54445    | 54445.0       |
| 10 | 63218  | 63218    | 63218.0       |
| 4  | 39891  | 39891    | 39891.0       |
| 1  | 46205  | 46205    | 46205.0       |
| 28 | 122391 | 122391   | 122391.0      |
| 8  | 64445  | 64445    | 64445.0       |
| 3  | 43525  | 43525    | 43525.0       |
| 23 | 113812 | 113812   | 113812.0      |
| 14 | 61111  | 61111    | 61111.0       |

```
In [27]: # y_test is series
# y_predictions is numpy array values
print(y_test.values[:5]) # float 5. means 5.0
print(y_predictions[:5])
```

```
[ 54445  63218  39891  46205 122391]
[ 54445.  63218.  39891.  46205. 122391.]
```

```
In [28]: # RMSE
# MSE
# MAE
# R-square

from sklearn.metrics import r2_score, mean_squared_error
```

```
In [29]: R2=r2_score(y_test,y_predictions)
MSE=mean_squared_error(y_test,y_predictions)
#MSE**(1/2)
RMSE=np.sqrt(MSE)
#accuracy_score(y_test,y_predictions) # it is a regression tech
print("R-sqaure:",R2)
print("MSE:",MSE)
print("RMSE:",RMSE)
```

R-sqaure: 1.0  
MSE: 6.470390569303684e-23  
RMSE: 8.043873798925294e-12

```
In [30]: s=0
for i in range(len(y_test)):
    v1=y_test.values[i]-y_predictions[i]
    v2=v1**2
    s=s+v2
print(s/len(y_test))
```

6.470390569303684e-23

```
In [31]: LR.coef_
print("The coeffiecnt of Years_of_experience is:",LR.coef_)
```

The coeffiecnt of Years\_of\_experience is: [1.]

```
In [32]: LR.intercept_
```

Out[32]: -1.4551915228366852e-11

```
In [33]: X_train.columns
```

Out[33]: Index(['Salary'], dtype='object')

```
In [34]: #Regression_equation=LR.intercept_+LR.coef_ * col namee
#Regression_equation

y=-1.45+1-.*Salary
```

Cell In[34], line 4

y=-1.45+1-.\*Salary

^

SyntaxError: invalid syntax

```
In [35]: from sklearn.feature_selection import VarianceThreshold
          vt=VarianceThreshold(threshold=0)
          # Threshold variance value
          # we want to drop the feaure based on threshold
          vt.fit(df)
```

```
Out[35]:
```

|                                |                   |
|--------------------------------|-------------------|
| ▼                              | VarianceThreshold |
| VarianceThreshold(threshold=0) |                   |

```
In [36]: dir(vt)
```



```
Out[36]: ['__abstractmethods__',
          '__annotations__',
          '__class__',
          '__delattr__',
          '__dict__',
          '__dir__',
          '__doc__',
          '__eq__',
          '__format__',
          '__ge__',
          '__getattribute__',
          '__getstate__',
          '__gt__',
          '__hash__',
          '__init__',
          '__init_subclass__',
          '__le__',
          '__lt__',
          '__module__',
          '__ne__',
          '__new__',
          '__reduce__',
          '__reduce_ex__',
          '__repr__',
          '__setattr__',
          '__setstate__',
          '__sizeof__',
          '__sklearn_clone__',
          '__str__',
          '__subclasshook__',
          '__weakref__',
          '_abc_impl',
          '_build_request_for_signature',
          '_check_feature_names',
          '_check_n_features',
          '_get_default_requests',
          '_get_metadata_request',
          '_get_param_names',
          '_get_support_mask',
          '_get_tags',
          '_more_tags',
          '_parameter_constraints',
          '_repr_html_',
          '_repr_html_inner',
          '_repr_mimebundle_',
          '_sklearn_auto_wrap_output_keys',
          '_transform',
          '_validate_data',
          '_validate_params',
          'feature_names_in_',
          'fit',
          'fit_transform',
          'get_feature_names_out',
          'get_metadata_routing',
          'get_params',
          'get_support',
          'inverse_transform',
          'n_features_in_',
          'set_output',
          'set_params',
          'threshold',
```

```
'transform',  
'variances_']
```

```
In [37]: vt.variances_  
# 300 is first column variance (T)  
# 1.25 is second column variance (T)  
# 30 is column variance (T)  
# 0 is fourth column variance (F)
```

```
Out[37]: array([7.78515556e+00, 8.46600000e+04])
```

```
In [38]: vt.get_support()
```

```
Out[38]: array([ True,  True])
```

```
In [39]: vt.get_params()  
# Hyper parameter  
# that we are providing inside the function
```

```
Out[39]: {'threshold': 0}
```

```
In [40]: vt.threshold
```

```
Out[40]: 0
```

```
In [41]: cols=vt.get_feature_names_out()  
# the above syntax gives the column names  
# These feature only we want include  
df[cols]
```

Out[41]:

|    | YearsExperience | Salary |
|----|-----------------|--------|
| 0  | 1.1             | 39343  |
| 1  | 1.3             | 46205  |
| 2  | 1.5             | 37731  |
| 3  | 2.0             | 43525  |
| 4  | 2.2             | 39891  |
| 5  | 2.9             | 56642  |
| 6  | 3.0             | 60150  |
| 7  | 3.2             | 54445  |
| 8  | 3.2             | 64445  |
| 9  | 3.7             | 57189  |
| 10 | 3.9             | 63218  |
| 11 | 4.0             | 55794  |
| 12 | 4.0             | 56957  |
| 13 | 4.1             | 57081  |
| 14 | 4.5             | 61111  |
| 15 | 4.9             | 67938  |
| 16 | 5.1             | 66029  |
| 17 | 5.3             | 83088  |
| 18 | 5.9             | 81363  |
| 19 | 6.0             | 93940  |
| 20 | 6.8             | 91738  |
| 21 | 7.1             | 98273  |
| 22 | 7.9             | 101302 |
| 23 | 8.2             | 113812 |
| 24 | 8.7             | 109431 |
| 25 | 9.0             | 105582 |
| 26 | 9.5             | 116969 |
| 27 | 9.6             | 112635 |
| 28 | 10.3            | 122391 |
| 29 | 10.5            | 121872 |

```
In [42]: path=r"C:\Users\kasho\OneDrive\Documents\Data Science\flask\Salary_Data - S
df=pd.read_csv(path)
df.head()
from sklearn.feature_selection import VarianceThreshold
vt=VarianceThreshold(threshold=0)
### Make sure before fitting the dataframe , do not include output column
X=df.drop('YearsExperience',axis=1)
# X it self a data frame
vt.fit(X)
vt.variances_
vt.get_support()
cols=vt.get_feature_names_out()
X[cols]
```

Out[42]:

|    | Salary |
|----|--------|
| 0  | 39343  |
| 1  | 46205  |
| 2  | 37731  |
| 3  | 43525  |
| 4  | 39891  |
| 5  | 56642  |
| 6  | 60150  |
| 7  | 54445  |
| 8  | 64445  |
| 9  | 57189  |
| 10 | 63218  |
| 11 | 55794  |
| 12 | 56957  |
| 13 | 57081  |
| 14 | 61111  |
| 15 | 67938  |
| 16 | 66029  |
| 17 | 83088  |
| 18 | 81363  |
| 19 | 93940  |
| 20 | 91738  |
| 21 | 98273  |
| 22 | 101302 |
| 23 | 113812 |
| 24 | 109431 |
| 25 | 105582 |
| 26 | 116969 |
| 27 | 112635 |
| 28 | 122391 |
| 29 | 121872 |

```
In [43]: from statsmodels.api import OLS
OLS(y_train,X_train).fit().summary()
```

Out[43]: OLS Regression Results

|                   |                  |                              |           |
|-------------------|------------------|------------------------------|-----------|
| Dep. Variable:    | Salary           | R-squared (uncentered):      | 1.000     |
| Model:            | OLS              | Adj. R-squared (uncentered): | 1.000     |
| Method:           | Least Squares    | F-statistic:                 | 3.694e+32 |
| Date:             | Tue, 16 Apr 2024 | Prob (F-statistic):          | 3.81e-314 |
| Time:             | 17:51:36         | Log-Likelihood:              | 488.13    |
| No. Observations: | 21               | AIC:                         | -974.3    |
| Df Residuals:     | 20               | BIC:                         | -973.2    |
| Df Model:         | 1                |                              |           |
| Covariance Type:  | nonrobust        |                              |           |

|        | coef   | std err | t        | P> t  | [0.025 | 0.975] |
|--------|--------|---------|----------|-------|--------|--------|
| Salary | 1.0000 | 5.2e-17 | 1.92e+16 | 0.000 | 1.000  | 1.000  |

|                |       |                   |       |
|----------------|-------|-------------------|-------|
| Omnibus:       | 3.403 | Durbin-Watson:    | 0.280 |
| Prob(Omnibus): | 0.182 | Jarque-Bera (JB): | 2.287 |
| Skew:          | 0.627 | Prob(JB):         | 0.319 |
| Kurtosis:      | 1.979 | Cond. No.         | 1.00  |

Notes:

[1] R<sup>2</sup> is computed without centering (uncentered) since the model does not contain a constant.

[2] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [46]: import pickle
pickle.dump(LR,
            open('YearsExperience_model.pkl','wb'))
```

```
In [47]: # Loading model to compare the result
model=pickle.load(open('YearsExperience_model.pkl','rb'))
model
```

Out[47]:

▼ LinearRegression

LinearRegression()

In [ ]:

In [ ]:

